

Applying Operating System Principles to SDN Controller Design

Matthew Monaco Oliver Michel Eric Keller

University of Colorado, Boulder

HotNets-XII
November 21, 2013



University of Colorado **Boulder**

NOX: Towards an Operating System for Networks

Natasha Gude
Nicira Networks

Ben Pfaff
Nicira Networks

Teemu Koponen
HIIT

Martín Casado
Nicira Networks

Scott Shenker
University of California,
Berkeley

Justin Pettit
Nicira Networks

Nick McKeown
Stanford University

is an editorial note submitted to
for this article's technique.

not been peer reviewed
will be posted there.

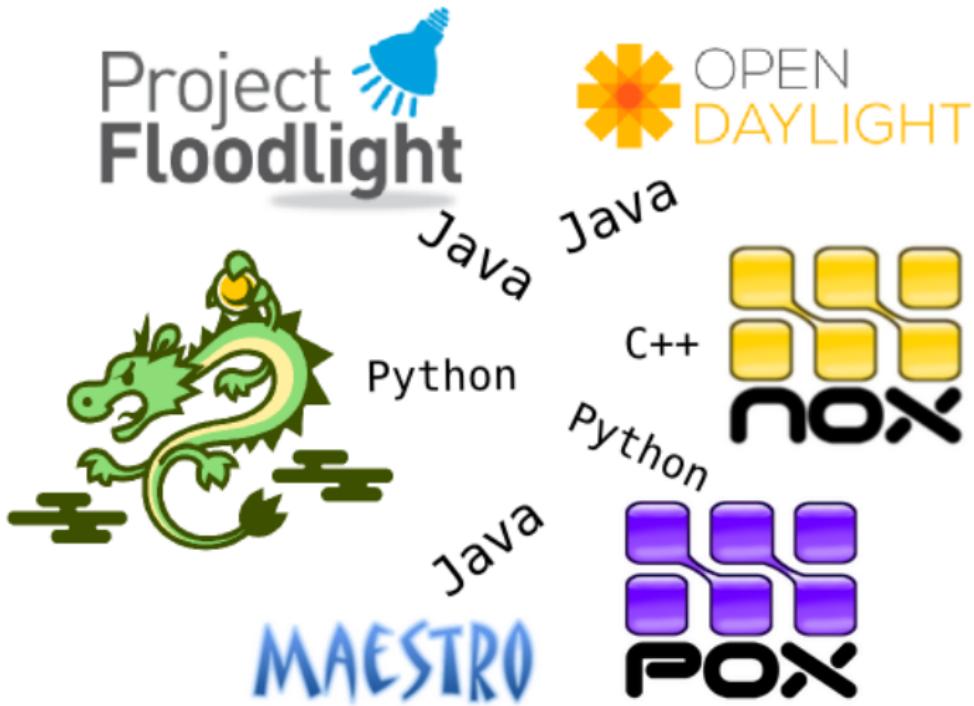
What we clearly need is an “operating system” for networks...

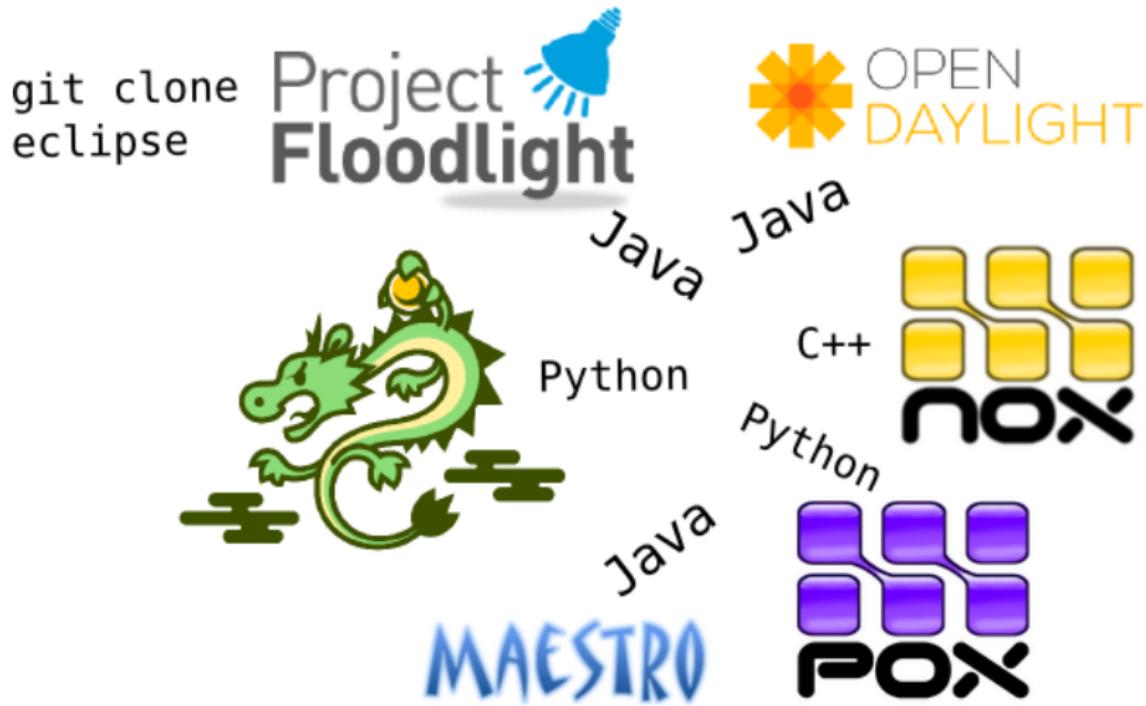




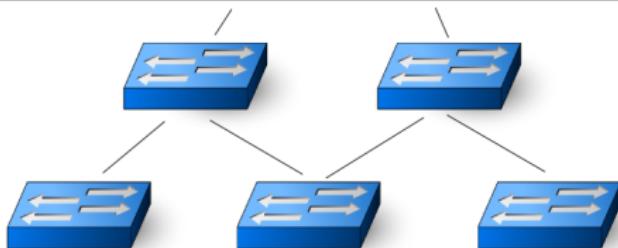
MAESTRO

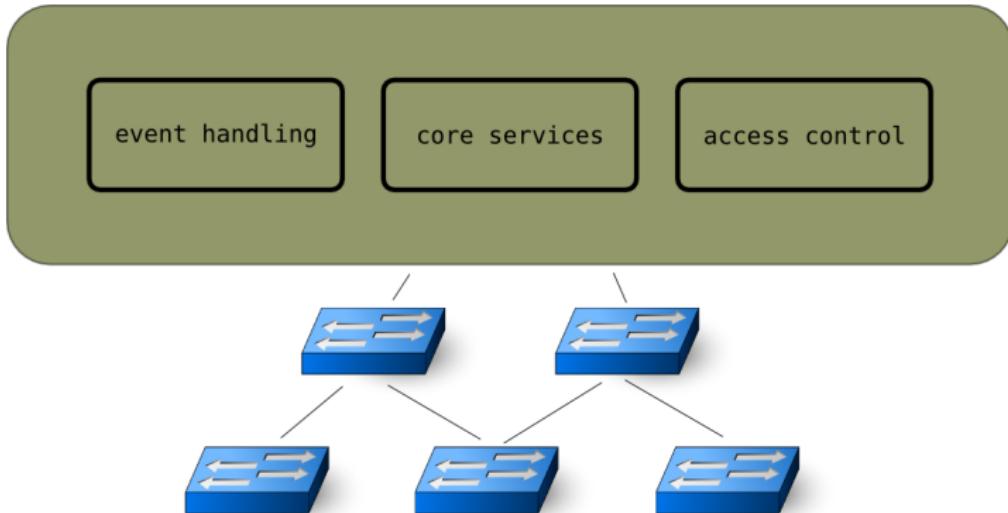






Monolithic





Traditional Operating Systems



Distributions



- Technically sound

Distributions



- Technically sound
- Distributions

Distributions



- Technically sound
- Distributions
 - » Core decisions
 - » Different Audiences
 - » Hackers vs. End Users

Distributions



- Technically sound
- Distributions
 - » Core decisions
 - » Different Audiences
 - » Hackers vs. End Users
- Networks

Distributions



- Technically sound
- Distributions
 - » Core decisions
 - » Different Audiences
 - » Hackers vs. End Users
- Networks
 - » Datacenters, Universities, Enterprises
 - » Small business, Home Users, Hobbyists

Software Projects



How do administrators and developers interact
with traditional OSes?



How do administrators and developers interact with traditional OSes?

\$



University of Colorado **Boulder**

How do administrators and developers interact with traditional OSes?

```
$
```

```
$ grep  
$ find  
$ sed  
$ sort  
$ cat
```



Detective Work

```
$ cd /var/log  
$ grep some_problem syslog
```



Performance Analysis

```
$ ps -A | sort -k3nr | head -10  
$ kill -TERM 12345
```



Simple Automation

```
#!/bin/bash

procs=(  
    ps -A | sort -k3nr | head -$1 | tr -s ' ' \  
    | cut -d' ' -f9  
)  
  
for p in "${procs[@]}"; do  
    printf "%d is misbehaving\n" "$p" >&2  
    kill -TERM "$p"  
    sleep 3  
    kill -KILL "$p"  
done
```



Large Applications

```
#include <stdlib.h>
#include <stdio.h>

int main(void)
{
    puts("Hello, World!");
    return EXIT_SUCCESS;
}
```

```
$ ./configure
$ make
$ sudo make install
```

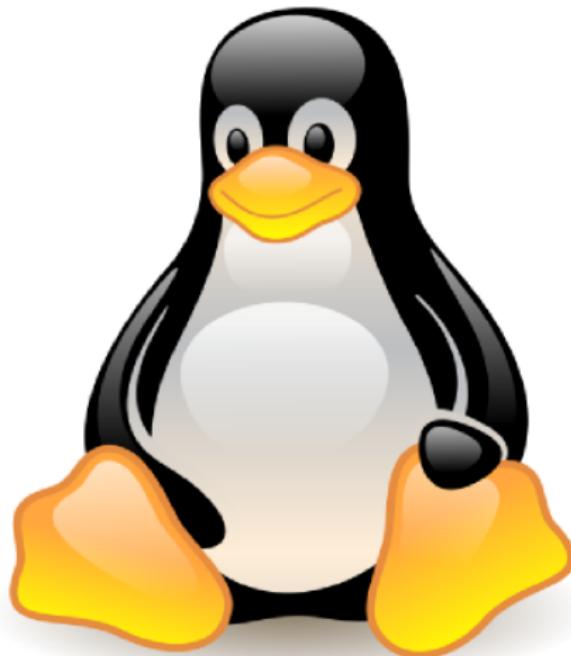


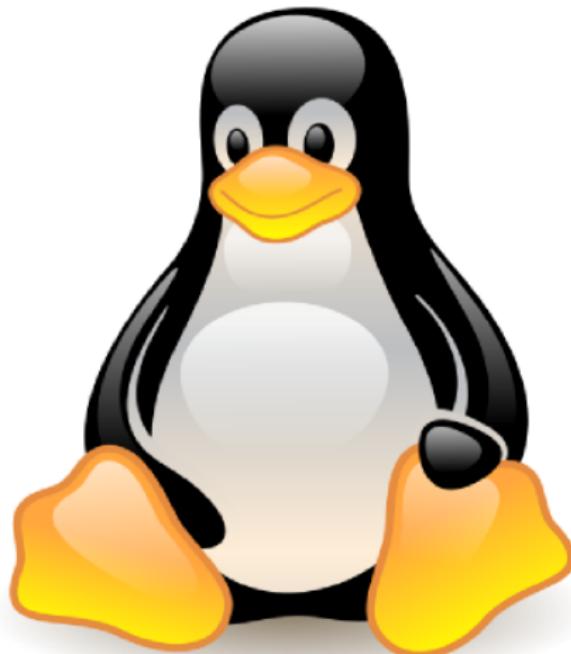
What can we learn from a traditional operating system?



What can we **use** from a traditional operating system?



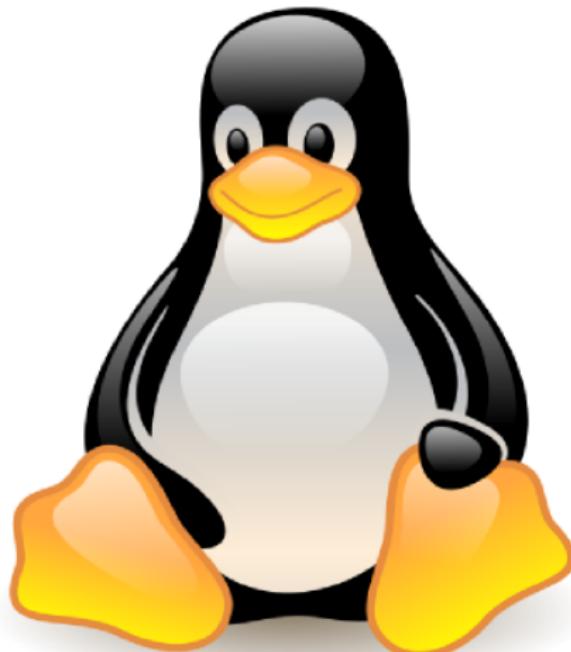




Integrate more tightly with
Linux

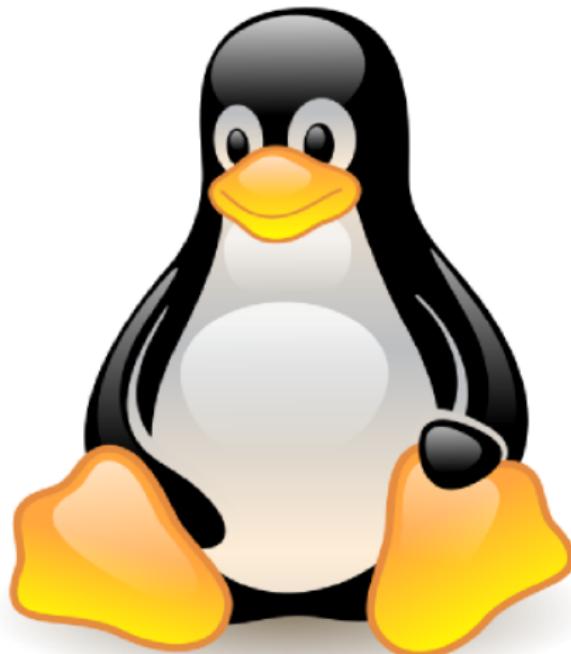


University of Colorado **Boulder**



Integrate more tightly with
Linux

Use off-the-shelf
technologies



Integrate more tightly with
Linux

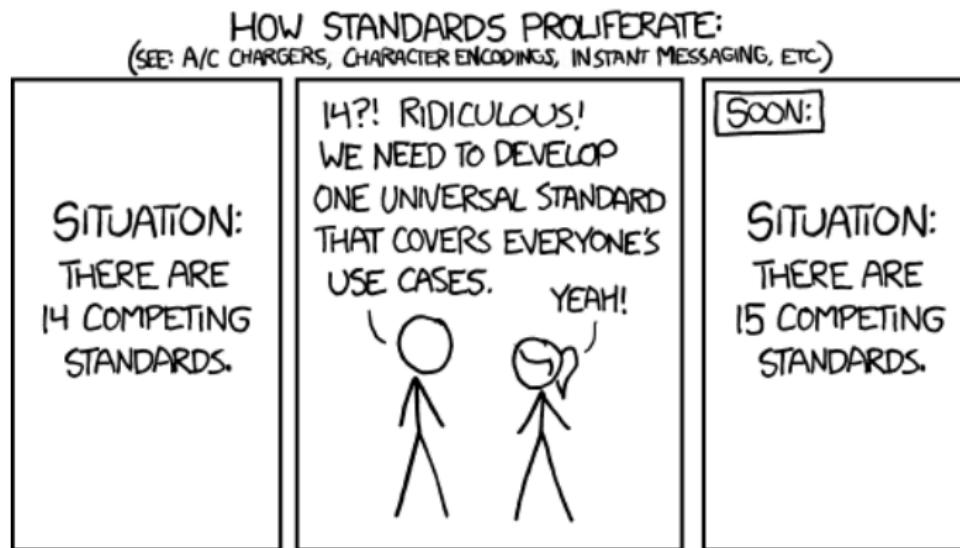
Use off-the-shelf
technologies

Encourage active
ecosystem

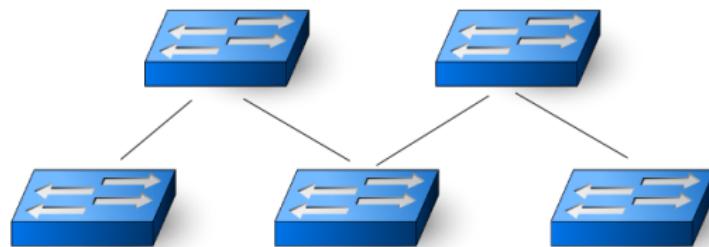
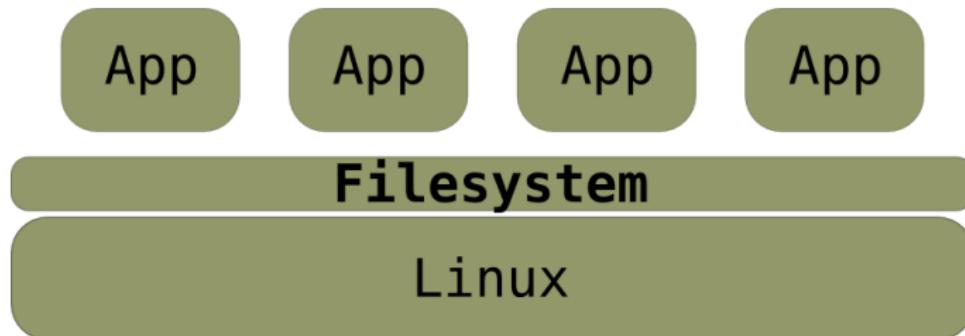
Yanc: Yet Another Network Controller



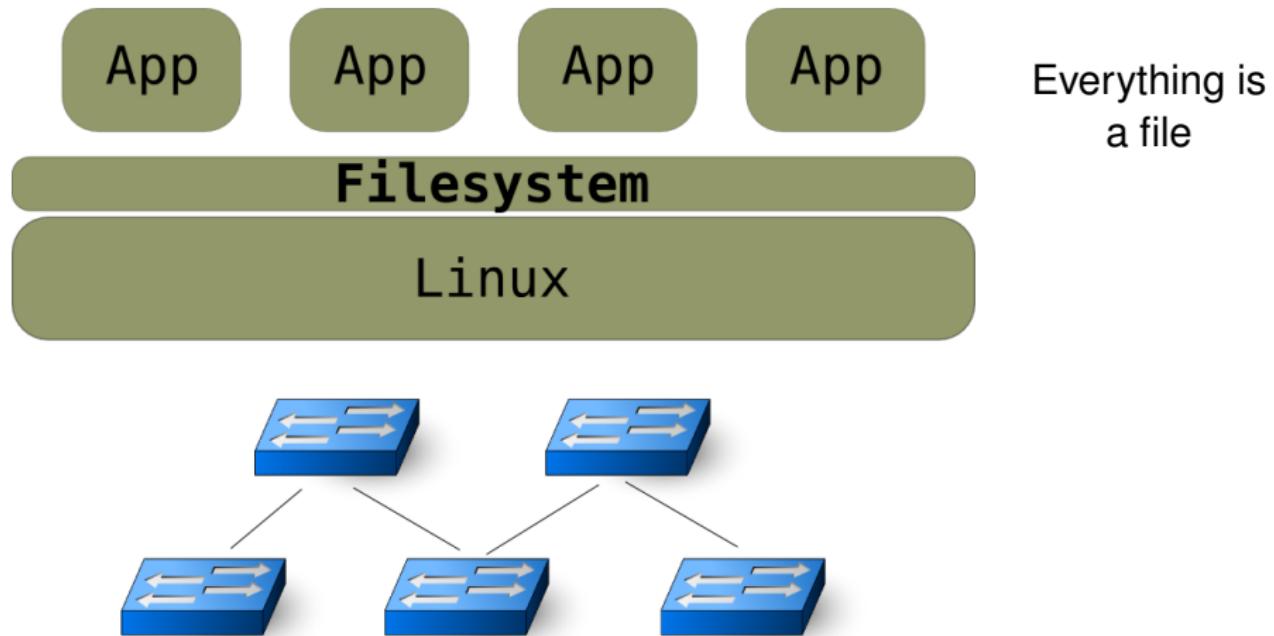
Obligatory XKCD



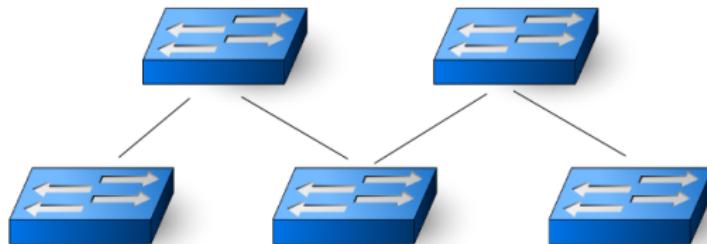
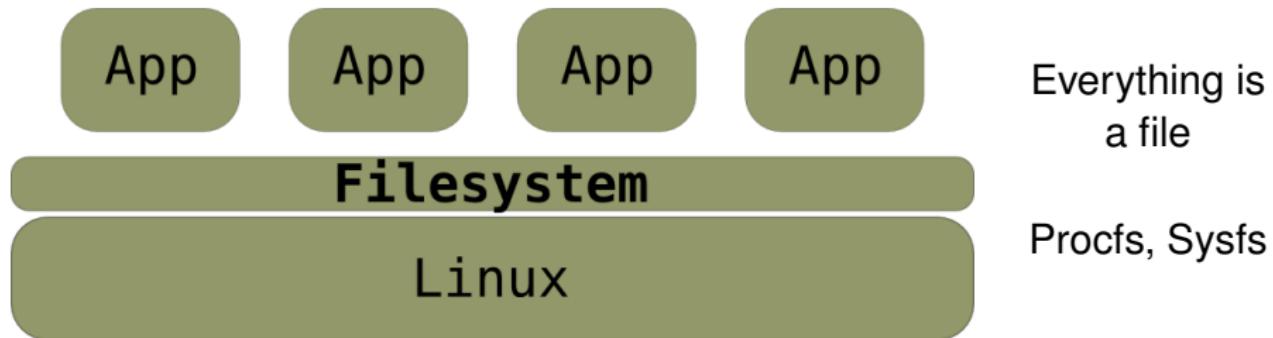
Yanc: Yet Another Network Controller



Yanc: Yet Another Network Controller



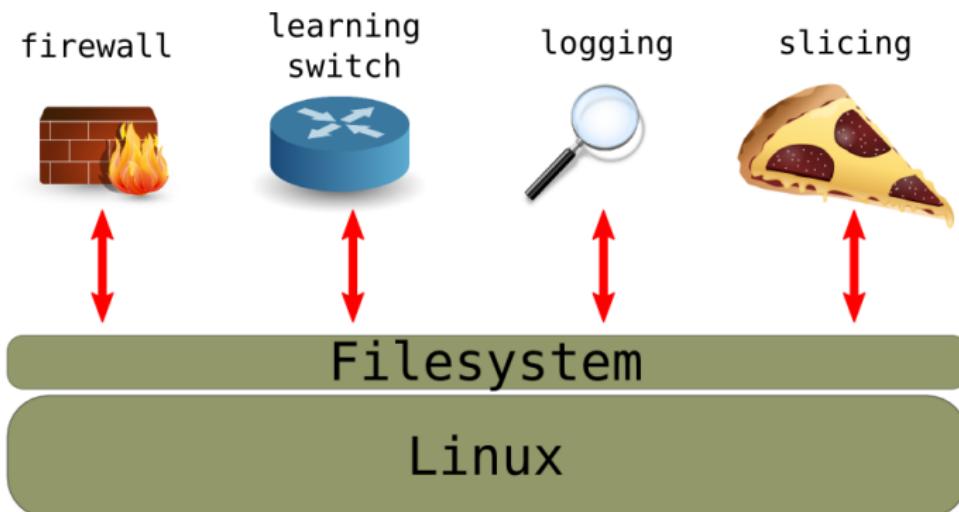
Yanc: Yet Another Network Controller



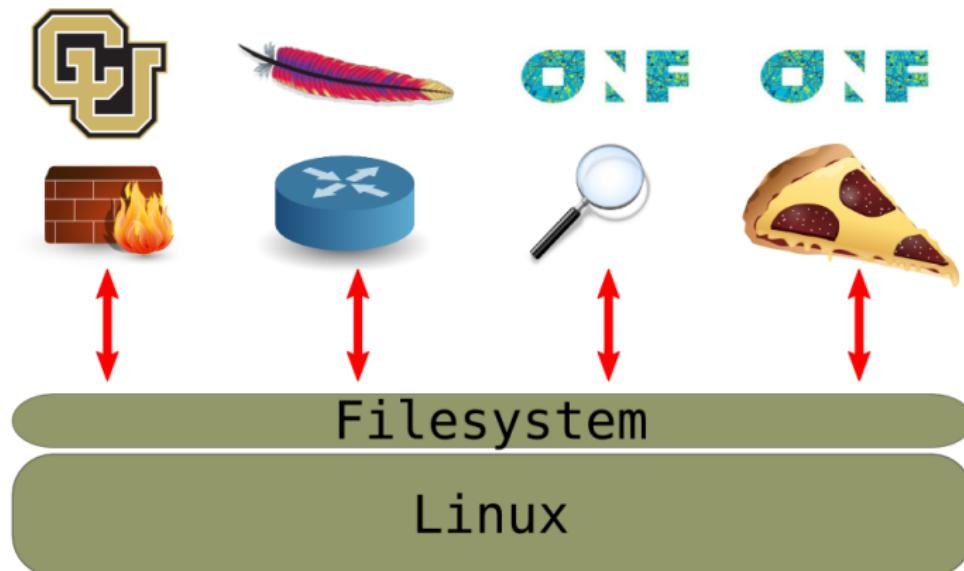
Why a filesystem?



Logically Distinct Applications



Independent Development



Independent Development

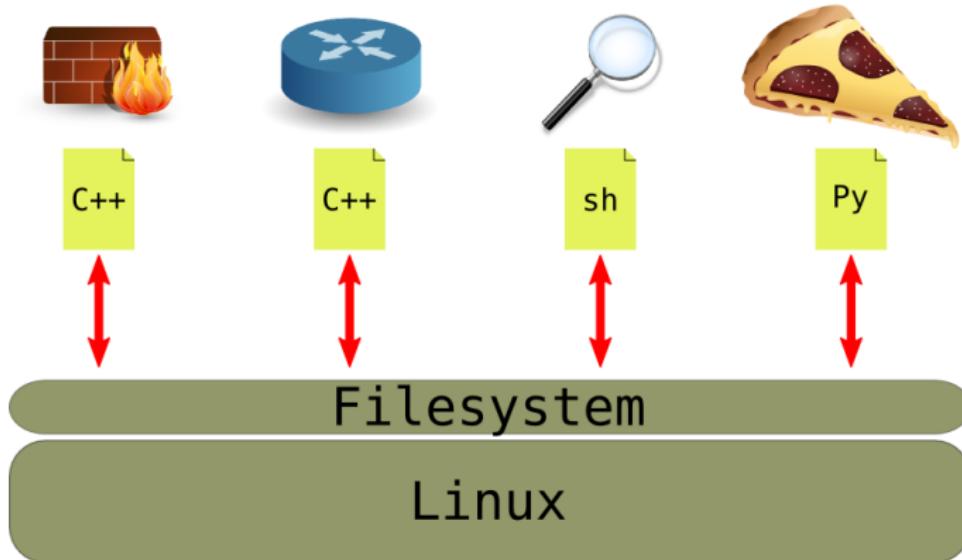
```
# apt-get install yanc-learning-switch
```

```
# apt-get install yanc-router
```

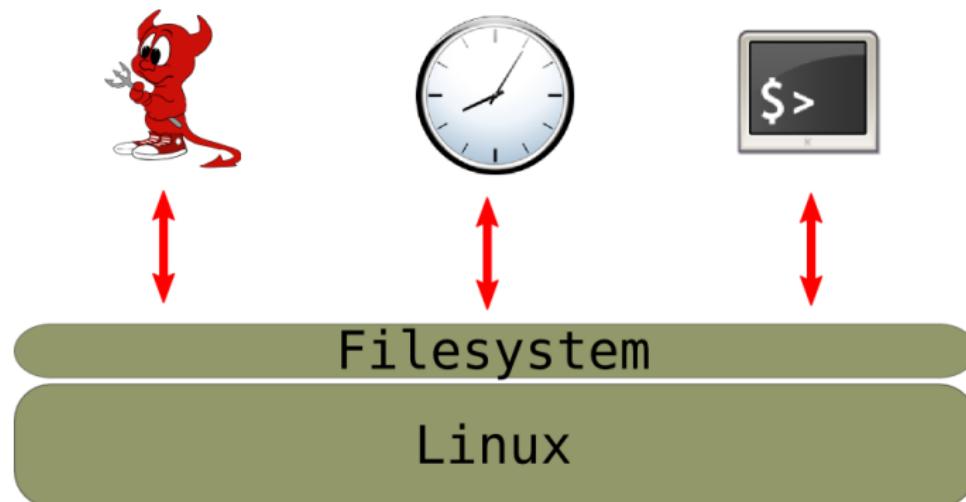
```
$ git clone git@github.com:mmonaco/yanc-firewall && cd  
    yanc-firewall  
$ make  
# make install
```



Any Programming Language



Any Form

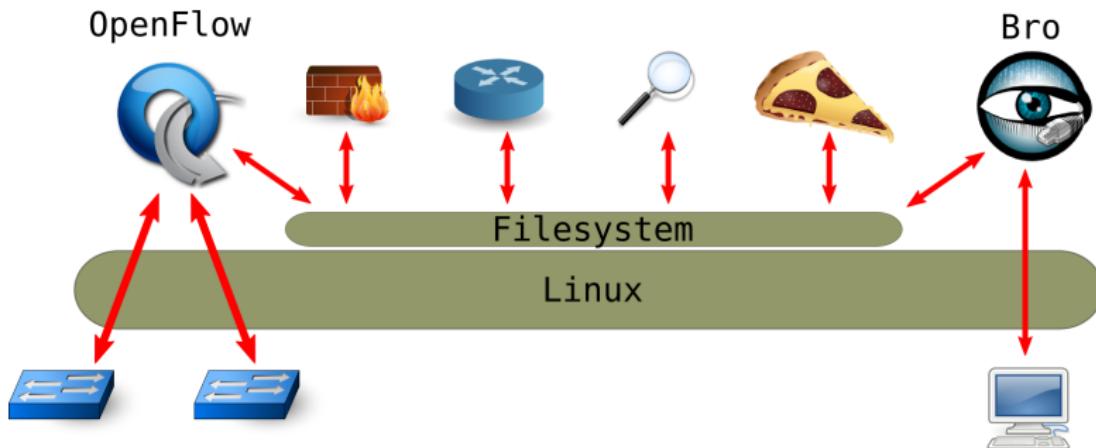


Other Technologies

- Inotify
- File permissions and ACLs
- Namespaces and CGroups
- Layered Filesystems



Decouple from Hardware



The Yanc Filesystem



Root Directory

```
/net
  └── hosts/
  └── switches/
      ├── sw1/
      └── sw2/
  └── views/
      ├── http/
      └── management-net
          ├── hosts/
          ├── switches/
          └── views/
```



Root Directory

```
/net
  └── hosts/
  └── switches/
      ├── sw1/
      └── sw2/
  └── views/
      ├── http/
      └── management-net
          ├── hosts/
          ├── switches/
          └── views/
```



Switch Directory

```
/net
└ switches
  └ 01:02:03:04:05:06
    ├── flows/
    ├── packet_in/
    ├── packet_out/
    ├── ports/
    ├── actions
    ├── capabilities
    ├── id
    └── n_buffers
```



Switch Directory

```
/net
└ switches
  └ 01:02:03:04:05:06
    ├── flows/
    ├── packet_in/
    ├── packet_out/
    ├── ports/
    ├── actions
    ├── capabilities
    ├── id
    └── n_buffers
```



Port Directory

```
/net
└ switches
  └ 01:02:03:04:05:06
    └ ports
      └ LOCAL
        └── config.port_down
        └── hw_addr
        └── peer -> /dev/null
        └── port_no
        └── stats.rx_bytes
        └── stats.rx_packets
```



Port Directory

```
/net
└ switches
  └ 01:02:03:04:05:06
    └ ports
      └ LOCAL
        ├── config.port_down
        ├── hw_addr
        ├── peer -> /dev/null
        ├── port_no
        ├── stats.rx_bytes
        └── stats.rx_packets
```



Flow Entry Directory

```
/net
└ switches
  └ 01:02:03:04:05:06
    └ flows
      └ arp_flow
        ┌── counters
        ┌── action.out
        ┌── match.dl_src
        ┌── match.dl_type
        ┌── priority
        ┌── timeout
        └── version
```



Flow Entry Directory

```
/net
└ switches
  └ 01:02:03:04:05:06
    └ flows
      └ arp_flow
        ┌── counters
        ┌── action.out
        ┌── match.dl_src
        ┌── match.dl_type
        ├── priority
        ┌── timeout
        └── version
```



Packet In Directory

```
/net
└ switches
  └ 01:02:03:04:05:06
    └ packet_ins
      └ in1
        └ buffer_id
        └ data
        └ data_len
        └ in_port
        └ reason
```



Packet In Directory

```
/net
└ switches
  └ 01:02:03:04:05:06
    └ packet_ins
      └ in1
        └ buffer_id
        └ data
        └ data_len
        └ in_port
        └ reason
```



Packet In Directory

```
/net
└ switches
  └ 01:02:03:04:05:06
    └ packet_ins
      └ in1
        └ buffer_id
        └ data
        └ data_len
        └ in_port
        └ reason
```



Packet Out Directory

```
/net
└ switches
    └ 01:02:03:04:05:06
        └ packet_outs
            └ out1
                └─ action.out_port
                └─ buffer_id
                └─ data
                └─ state
```



Packet Out Directory

```
/net
└ switches
    └ 01:02:03:04:05:06
        └ packet_outs
            └ out1
                └─ action.out_port
                └─ buffer_id
                └─ data
                └─ state
```



Using the Filesystem



```
$ echo 1 > port_1.port_down
```



```
$ echo 1 > port_1.port_down
```

```
$ cd switches/00:01:02:0a:0b:0c/flows
$ mkdir my_flow_entry
$ ls -1 my_flow_entry
counters
priority
timeout
version
```



```
$ find /net -name tp.dst -exec grep 22 {} +
```



```
$ find /net -name tp.dst -exec grep 22 {} +
```

```
#!/bin/bash
flowdir=/net/switches/"$1"/flows/"$2"
mkdir "$flowdir"
echo ff:ff:ff:ff:ff:ff > "$flowdir"/match.dl_dst
echo 0x0806 > "$flowdir"/match.dl_type
echo FLOOD > "$flowdir"/action.out
```



Libraries

```
#!/usr/bin/env python3

def new_switch(id, n_tables=1):
    pass

def write_flow(switch, matches=[], actions=[]):
    pass
```

```
#ifndef _YANC_H_
#define _YANC_H_

int new_switch(uint64_t, uint8_t);
int write_flow(const char* path, match_t*, action_t*);

#endif/*_YANC_H_*/
```



- Filesystem - C, FUSE



- Filesystem - C, FUSE
- OpenFlow - C++



- Filesystem - C, FUSE
- OpenFlow - C++
- Discovery - Python



- Filesystem - C, FUSE
- OpenFlow - C++
- Discovery - Python
- Static Flow Pusher - Bash



Ongoing Work



Composition

Distribution

Applications



Composition

- Independent development



Composition

- Independent development
- Event handling



Composition

- Independent development
- Event handling
- Synthesize configuration



Composition

- Independent development
- Event handling
- Synthesize configuration
- /usr/share/yanc.d/
- /etc/yanc.d/



Composition

- Independent development
- Event handling
- Synthesize configuration
- /usr/share/yanc.d/
- /etc/yanc.d/
- Ship with configuration



Composition

- Independent development
- Event handling
- Synthesize configuration
- /usr/share/yanc.d/
- /etc/yanc.d/
- Ship with configuration
- Administrator controlled



Composition

```
# /etc/yang.d/  
  
<event> slice  
<event> firewall  
<event> router
```



Distribution

Latency

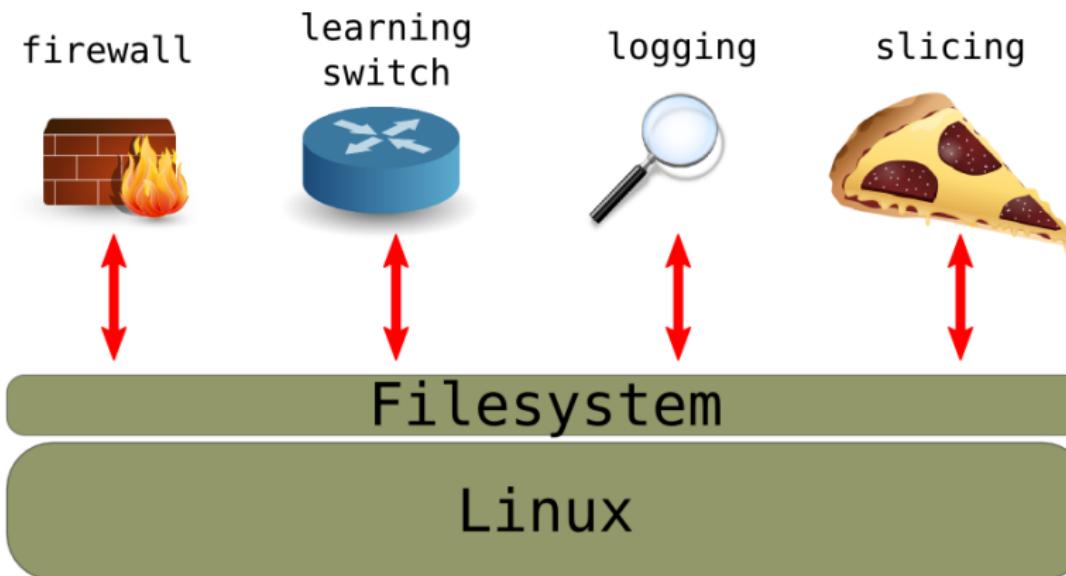
Fault Tolerance

Administration

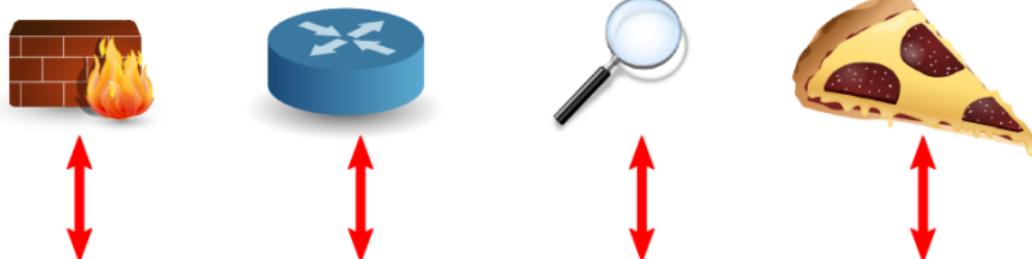
Different Requirements



Distribution



Distribution

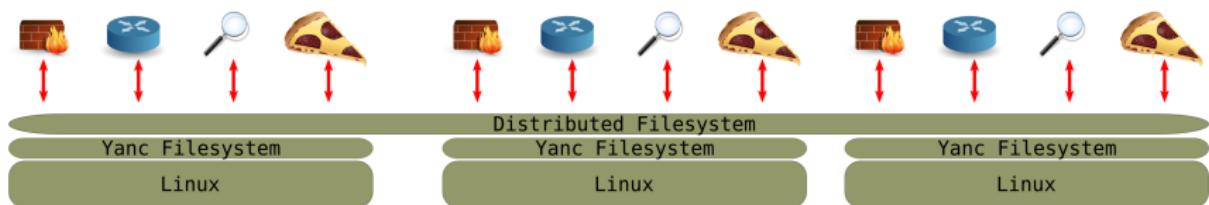


Distributed Filesystem

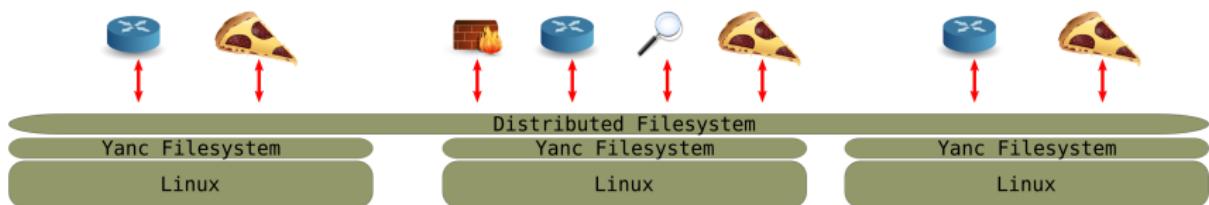
Yanc Filesystem

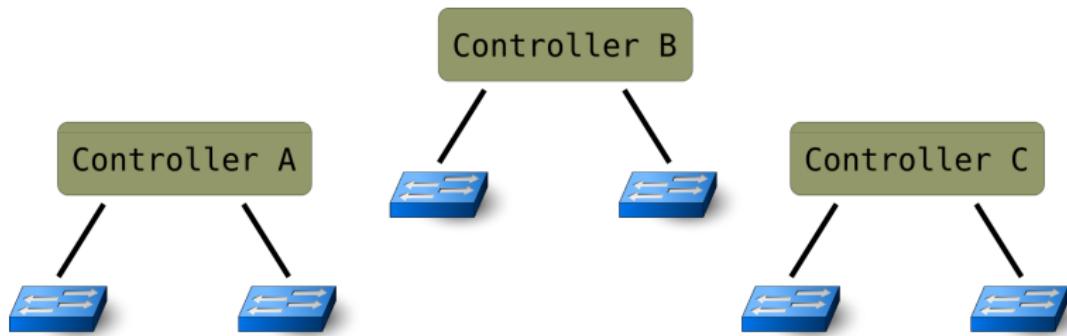
Linux

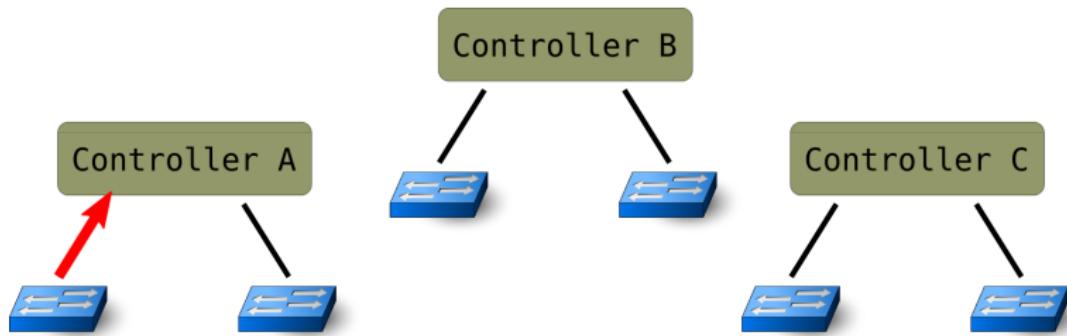
Distribution

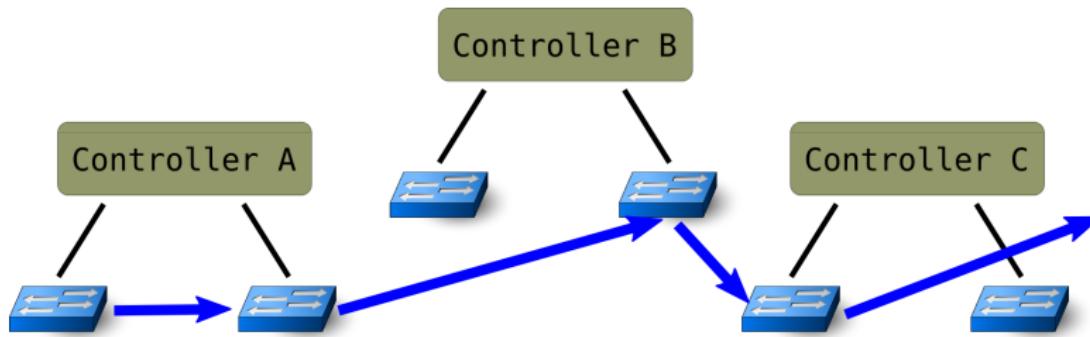


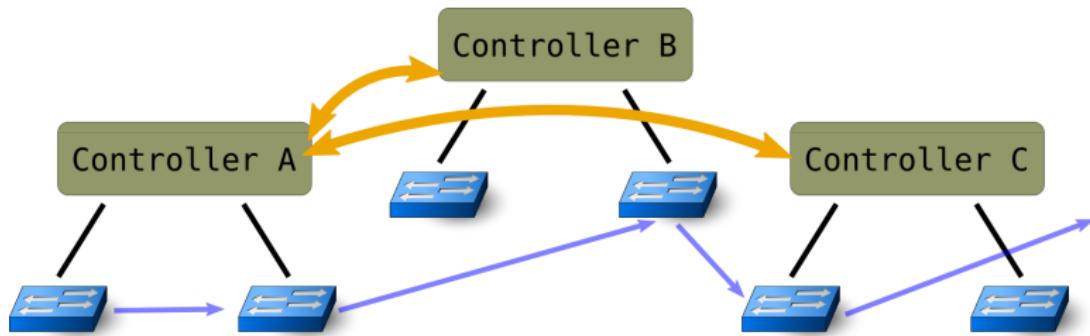
Distribution

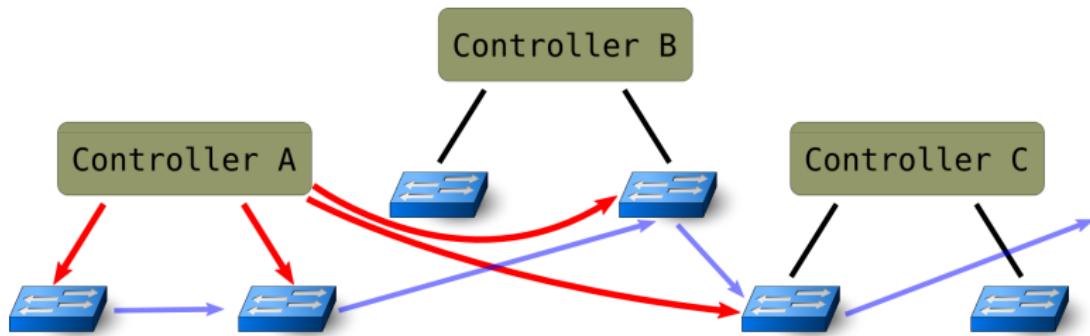


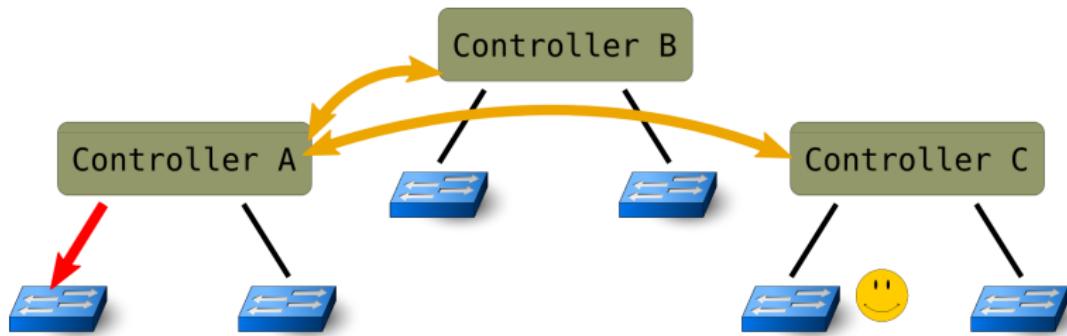












Defined a Universal Interface for Network Control



Defined a Universal Interface for Network Control

Allow Control Applications as Separate Processes



Defined a Universal Interface for Network Control

Allow Control Applications as Separate Processes

Implemented Interface



Defined a Universal Interface for Network Control

Allow Control Applications as Separate Processes

Implemented Interface

Leveraging Existing OS Technologies



Defined a Universal Interface for Network Control

Allow Control Applications as Separate Processes

Implemented Interface

Leveraging Existing OS Technologies

Built Functional Applications on Top of Interface



Thank You!



University of Colorado **Boulder**